

# Reinforcement Learning 2

Complications & approximations

---

Andrew Lampinen



Psych 209, Winter 2018

# Introduction

---

# Plan for this lecture

Talk about all the stuff that makes it messy:




- Infinite spaces, approximation & generalization.
- Correlations & replay.
- Exploration & on/off-policy learning.
- Hierarchies & plans.
- Other types of feedback.
- Unobservables & POMDPs, different assumptions and other approaches.

# Function approximation

---

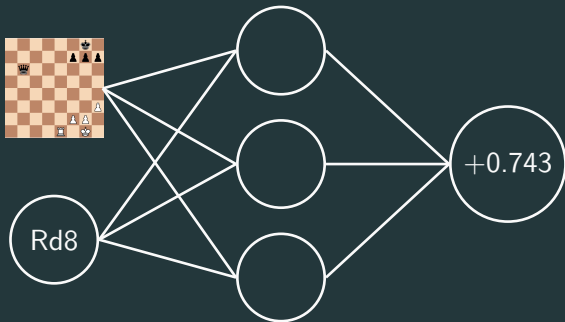
# Similarities & dissimilarities

- Wheels of car can be at infinitely many angles, but is 0.55 radians really that different from 0.56?
- More state, action pairs in chess than atoms in observable universe. However:



## Approximating the Q-table

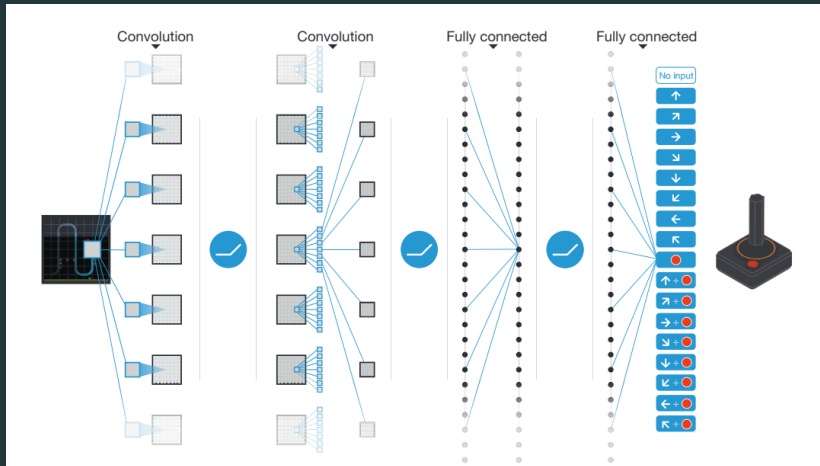
- Previously, the Q-table was a lookup function. Let's replace this with some other function that maps states and actions to Q-values.



- Loss:

$$L = \left( r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a'; \theta_i^-) - Q(s_t, a_t; \theta_i) \right)^2$$

# Playing atari games




# Replay

---



## Fix one problem and ...

- Generalization comes at the expense of cross talk and catastrophic forgetting. Will a self-driving car learning to parallel park forget how to drive on the freeway?
- CLS is a theory of how humans and animals avoid this:



## Replay buffers

The Atari game-playing paper addressed this with **experience replay**:

- Whenever we experience a new  $(s_t, a_t, r_{t+1}, s_{t+1})$  tuple, stick it in a buffer.
- instead of

$$\Delta Q^\pi(s_t, a_t) = \alpha \underbrace{\left( \left[ r_{t+1} + \max_{a'} \gamma Q^\pi(s_{t+1}, a') \right] - Q(s_t, a_t) \right)}_{\text{prediction error!}}$$


- we sample a random experience from our buffer and update with that:  $k \sim \text{Unif}(\text{replay buffer indices})$

$$\Delta Q^\pi(s_k, a_k) = \alpha \underbrace{\left( \left[ r_{k+1} + \max_{a'} \gamma Q^\pi(s_{k+1}, a') \right] - Q(s_k, a_k) \right)}_{\text{prediction error!}}$$

## Exploration and exploitation and on/off policy

---

# The need to explore



- How do I trade off between **exploiting** the best path to lunch I've found so far and **exploring** my other options?
- In other words, how do we incorporate exploration into our policy?
- This is important – remember convergence guarantees required *every* state, action pair to be visited “frequently.”

# The exploration-exploitation trade-off

There is a *fundamental* trade-off between exploring and exploiting:

- Exploring wastes time trying things I'm pretty sure aren't good.
- Exploiting risks missing a great opportunity.
- We have to find some balance between these that results in good payoffs but still makes sure we don't miss too much.


## $\epsilon$ -greedy

- One simple technique is just to choose actions randomly some small fraction  $\epsilon$  of the time.
- The rest of the time, take the maximum  $Q$  action.
- We call this  $\epsilon$ -greedy.
- We can also do clever things like *anneal*  $\epsilon$  over training.
  - Act mostly randomly and explore a lot early in training.
  - Act mostly greedily and exploit a lot late in training/in testing.
- (There are other possibilities, e.g. using a softmax over  $Q$  values.)

# Exploit during testing or keep exploring?

The point about testing vs. training is a little tricky...

- Do chess players stop learning when they're playing for the world championship?
- A self-driving car can't just be trained and set free – destinations, roads, and laws are all evolving.
- What if my direct path to lunch was blocked by a building that was torn down?



## Non-optimality of Q-learning under exploration

- But Q-learning fundamentally assumes that we will be behaving completely greedily during training! Why?
- It's built into our computation of the Q-update:

$$\Delta Q^\pi(s_t, a_t) = \alpha \left( \left[ r_{t+1} + \max_{a'} \gamma Q^\pi(s_{t+1}, a') \right] - Q(s_t, a_t) \right)$$

- If we want to behave optimally with a policy  $\pi$  that explores, we have to incorporate this policy into the learning rule:

$$\Delta Q^\pi(s_t, a_t) \stackrel{?}{=} \alpha \left( \left[ r_{t+1} + \sum_{a'} p(a' | s_t, a_t, s_{t+1}, \pi) \gamma Q^\pi(s_{t+1}, a') \right] - Q(s_t, a_t) \right)$$



# SARSA

- One algorithm for doing this is SARSA. It works almost exactly like Q learning, except we store  $(s_t, a_t, r_{t+1}, s_{t+1}, a_{t+1})$
- and change our update from

$$\Delta Q^\pi(s_t, a_t) = \alpha \left( \left[ r_{t+1} + \max_{a'} \gamma Q^\pi(s_{t+1}, a') \right] - Q(s_t, a_t) \right)$$

to

$$\Delta Q^\pi(s_t, a_t) = \alpha ([r_{t+1} + \gamma Q^\pi(s_{t+1}, a_{t+1})] - Q(s_t, a_t))$$

- Since our actions in the state are distributed according to  $\pi$ , in expectation our Q-value updates will be as well.

## On/off-policy

This highlights an important, but somewhat orthogonal distinction:

- Do we want to learn **on-policy**, that is, using the same policy we will be using during evaluation (like SARSA)?
- Or do we want to learn **off-policy**, that is, using a different policy during training than during evaluation (like Q-learning)?


There are tradeoffs!

- **On-policy:** Can be faster, can be more stable.
- **Off-policy:** Can learn from anything, even totally random play. This makes incorporating replay or changing  $\epsilon$  for more early exploration easier.

# Hierarchies & plans

---

## Plans = hierarchies of actions



## Hierarchies of Q-learning?

- You could think of doing Q-learning at each level, where state includes states of the above levels (including their goals and current higher-level actions, etc).
- But how do you figure out what the higher level actions should be?
- Still an open problem (though there is work on it), and potentially a good project is to think about how some aspect of behavior could be explained this way, and what assumptions you would need to make it work!

## Other types of feedback

---

# Auxiliary tasks

- Actually, we might get other forms of feedback... Can we incorporate these?
- Short answer: yes.
- May help with picking up on longer-term structure.
- Or with building relevant prior knowledge.



## Other approaches

---



## Other approaches

Many other approaches:

- Policy gradient methods.
- Actor-critic methods.
- Many, many variations.

And under different assumptions:

- What if we only observe part of the state? (POMDPs.)
- What if we actually build models of the world and plan over these? (Model-based RL, as opposed to Model-free.)
- ... Or something in between? (Successor-representation, imagination.)
- And much more.

## Wrapping up

---

# Summary

- We have an **agent** that takes **actions** in **states**.
- These actions are chosen according to a **policy**, sometimes derived from some sort of **value function**.
- We evaluated on **expected discounted rewards**.
- The **discount** reflects how future-oriented we are.
- $Q^\pi(s, a)$  is the expected return of  $a$  in  $s$  under  $\pi$ .
- We can learn  $Q$  values by **TD learning** (surprise).
- Can approximate  $Q$  using deep learning (for generalization).
- There's a trade-off between **exploring** and **exploiting**!
- ... and there's way more to RL than will fit on one slide or in one course.